



RAG Flow Explanation

RAG flow



Then when the user asks a question:



Simple example


User asks: "What is the company's leave policy?"

The question is converted into a **query vector**.

The vector database compares that query vector with stored vectors using a similarity method such as cosine similarity.

It might retrieve: *"Employees receive 20 days of annual leave per year..."*

That retrieved text becomes **context**.



One important point

The vector itself is **NOT** usually sent to the LLM as the context.

Query → **vector** → **similarity search** → **relevant text/chunks** → **context** → **LLM** → **answer**

So your understanding is essentially **100% correct**.